

Safety-Monitored Edge Perception for ADAS: Connecting Runtime Monitoring Evidence to a Bounded Safety Argument

Pruthi Parikh — draft v0.9 (Week 9). IEEE-style structure; PDF build in Week 10. All numbers generated by scripts/build_report_assets.py from results/* (see paper/tables.md); no hand-typed metrics.

Abstract

Deep-learning perception is the least certifiable component of an ADAS stack: it fails silently, out of its training distribution, without an error signal. This work builds a camera-only pedestrian/vehicle/cyclist detector for an AEB-style function and — rather than maximizing accuracy — asks the safety-engineering question: *when should this stack not be trusted, and can that be detected at runtime within a real-time budget?* We train YOLOv8n on KITTI (val mAP50 0.859), deploy via TensorRT FP16 (mAP50 delta -0.0024 , pure-inference 2.6 ms), calibrate confidences by temperature scaling (ECE 0.081 $\rightarrow$ 0.039), and gate the output with a frame-level monitor: a max-confidence OOD score with Q95/Q99 thresholds frozen from validation data drives a three-state machine (NOMINAL/DEGRADED/FAIL_SAFE_REQUEST) triggering only on sustained breaches. On BDD100K shift slices the score separates night from in-distribution data at AUROC 0.982; at the frozen Q95 threshold the monitor retains 94.8% of in-distribution frames while accepting only 6% of night frames. Under synthetic fault injection the monitor flags 99% of frames as severe fog collapses mAP50 from 0.847 to 0.071, with the full perception+monitor loop at p95 17.2 ms against a 40 ms budget. Equally important, the campaign quantifies a blind spot: severe low-light removes 19% of mAP50 while the monitor stays 97.7% NOMINAL. We package both results — capability and documented residual — into a GSN safety case, SOTIF argument, and ISO/PAS 8800-aligned traceability from hazards (STPA/HARA) to per-requirement evidence. The contribution is methodological: an end-to-end, reproducible template connecting edge-ML performance artifacts to safety-engineering evidence, with claims bounded to a defined ODD and honest residual risks.

I. Introduction

Advanced driver-assistance functions such as automatic emergency braking (AEB) place a neural detector in a safety-critical control loop. Two properties make this uncomfortable for classical functional-safety practice: the detector's failure modes are input-dependent rather than component-fault-driven, and its confidence outputs are not calibrated probabilities. ISO 21448 (SOTIF) frames this as *performance limitations under triggering conditions*; ISO/PAS 8800 extends the lifecycle view to AI-specific evidence. What is often missing in portfolio and research prototypes is the connective tissue: metrics exist, safety documents exist, but nothing traces a specific AUROC number to a specific hazard.

The work grew out of two earlier threads: encountering YOLO during an IoT seminar research project, which made edge-deployed real-time detection concrete, and an ECG digital-twin project, which framed safety as an engineering obligation to reason about when a model should not be trusted and to make that failure observable. This project is their deliberate intersection — an edge detector treated as a component whose untrustworthy regions must be surfaced, not smoothed over.

This project builds that connective tissue for a deliberately narrow item: camera-only 2D detection of pedestrians, vehicles, and cyclists within a KITTI-like operational design domain (ODD), supervised by a runtime monitor that requests degraded or fail-safe handling when perception evidence degrades. Contributions:

1. **A monitored perception MVP on edge hardware** (RTX 3050 Ti, 4 GB): YOLOv8n + TensorRT FP16, calibrated confidence, frame-level OOD score, three-state gating, per-frame evidence logging — p95 17.2 ms full loop.
2. **Threshold and evaluation hygiene as a first-class design rule**: thresholds from validation quantiles only; test and shift slices report-only; seeds and split hashes committed.
3. **A two-campaign triggering-condition exploration** (real BDD100K night/rain/fog slices; deterministic synthetic corruptions) showing where frame-level max-confidence monitoring works — and quantifying where it does not.

4. **A bounded safety argument:** STPA/HARA-derived requirements SR-01..SR-06, each verified by scripted evidence; GSN case with an explicit residual-risk node; SOTIF classification moving the discovered low-light blind spot from unknown-unsafe to known-unsafe.

We claim alignment with safety-engineering practice, not compliance or certified safety.

II. Related Work

Runtime monitoring of DNN perception. Max-softmax/max-confidence baselines (Hendrycks & Gimpel), energy-based OOD scores (Liu et al.), and Mahalanobis feature-space detectors (Lee et al.) define the standard spectrum from output-level to representation-level monitors. We deploy the simplest (max-confidence) after empirically finding our post-NMS energy proxy no better (mean OOD AUROC 0.913 vs 0.928), and carry feature-space methods as future work — an explicit scope decision recorded before evaluation on test data.

Calibration. Temperature scaling (Guo et al.) remains the standard post-hoc method; we apply a scalar T in logit space to detection confidences with correctness defined by class-matched IoU ≥ 0.5 greedy assignment.

Safety engineering for ML perception. STPA (Leveson) supplies unsafe-control-action analysis suited to software-intensive systems; SOTIF (ISO 21448) frames triggering conditions; GSN (SCSC) structures the argument; ISO/PAS 8800 names AI-lifecycle evidence obligations. Our use is educational-depth alignment on a real evidence base.

Robustness benchmarks. Corruption benchmarks (Hendrycks & Dietterich; KITTI-C/CODA variants) motivate our synthetic campaign; we use lightweight deterministic OpenCV corruptions rather than benchmark suites to keep the dependency stack frozen and every image bit-reproducible from (frame-id, corruption, severity).

III. System Design

Item and ODD. Forward camera, KITTI-like mounting; three classes; daylight, clear/lightly overcast, KITTI-style road scenes. Night, adverse weather, and degraded imaging are **outside** the ODD — they are precisely what the monitor targets. Downstream planner/AEB and actuators are simulated boundaries; assumption A-03 (planner honors monitor states) is untested by design and carried as the largest interface risk.

Architecture. camera $\rightarrow$ YOLOv8n (PyTorch | TensorRT FP16) $\rightarrow$ runtime monitor $\rightarrow$ {gated detections, monitor state} $\rightarrow$ planner boundary, with a per-frame machine-readable log as the evidence spine. Module boundaries in code mirror this: src/dataset, src/monitor/{calibration,scoring,thresholds,state_machine,runtime}, scripts/* as the only metric-producing entry points.

Safety derivation. HARA-lite works one hazardous event fully (undetected pedestrian in ego path, urban daytime: S3/E4/C3 $\rightarrow$ ASIL D as an educational classification) yielding safety goal SG-01; STPA over two control actions (gated detections, monitor state) yields 9 unsafe control actions and 9 causal scenarios, which decompose SG-01 into six measurable requirements: calibrated confidence (SR-01), runtime OOD scoring (SR-02), sustained-breach degraded entry (SR-03), fail-safe request (SR-04), per-frame logging (SR-05), 40 ms p95 latency budget (SR-06).

IV. Monitor Design

Calibration (SR-01). kitti-val (1122 frames) splits 50/50 (seed 42) into calibration-fit and calibration-report subsets. Detections at $\text{conf} \geq 0.05$ match GT greedily (same class, $\text{IoU} \geq 0.5$, one-to-one, confidence-descending). A scalar temperature minimizing binary NLL on the fit subset (golden-section search) gives $T = 0.600$ — the model is **underconfident**; ECE on the report subset improves $0.0812 \rightarrow 0.0390$ (15 equal-width bins).

OOD score (SR-02). Frame-level $\text{max_conf_score} = 1 - \max(\text{conf})$; empty frames score 1.0 (no evidence of in-distribution input — empirically the dominant failure signature under severe degradation). A post-NMS energy-style proxy ($-\log\text{sumexp}$ over top-10 confidence logits) is evaluated as the committed comparison and found not materially better; the simpler score is selected by pre-registered policy (energy must win by ≥ 0.01 mean AUROC).

Thresholds (SR-03/04). Q95 = 0.1862 and Q99 = 0.3728: quantiles of the ID score distribution on the calibration-report subset only. BDD and kitti-test never touch threshold selection.

Gating. Sustained-breach state machine: 3 consecutive Q95 breaches enter DEGRADED; 2 consecutive Q99 breaches or 10 unrecovered degraded frames enter FAIL_SAFE_REQUEST; 5 consecutive clean frames recover one level; breach is strict (score > threshold). Single-frame spikes cannot transition — verified deterministically.

V. Experimental Evaluation

All experiments: seed 42, committed split hashes, commands in docs/experiment_log.md (EXP-001..EXP-011); tables in paper/tables.md.

A. Baseline (Table 1). YOLOv8n, 50 epochs: val mAP50 0.8588 (pedestrian 0.749, vehicle 0.948, cyclist 0.880), mAP50-95 0.556. TensorRT FP16: mAP50 0.8564 (Δ -0.0024), pure engine inference 2.6 ms/img; end-to-end p50 16.9 ms.

B. Calibration (Table 2). ECE 0.0812 $\rightarrow$ 0.0390 on the disjoint report subset; reliability diagrams confirm the pre-scaling accuracy-above-diagonal signature of underconfidence.

C. OOD separation (Table 3). Against the ID report subset, max-conf AUROC/FPR@95: night 0.982/0.112; rain 0.876/0.537; fog 0.926/0.419 (n=13); clear-day transfer control 0.758/0.624. The transfer control matters: even in-ODD-like BDD imagery is detectably shifted from KITTI, so thresholds cannot transfer across camera domains without re-derivation.

D. Risk-coverage (Table 4). At frozen Q95 the monitor keeps 94.8% of ID frames (accepted-frame mAP50 0.854 vs 0.850 at full coverage) while accepting only 6.0% of night, 36.0% of rain, 23.1% of fog, and 68.0% of clear-day-control frames. The monitor's value is out-of-ODD rejection, not ID accuracy gain — accepted-frame mAP is nearly flat in coverage, and we report that honestly.

E. Runtime and latency (Table 5). Full loop (predict + score + state machine + log row): TRT FP16 p50 14.6 / p95 17.2 ms (67.9 FPS); PyTorch p50 15.7 / p95 17.5 ms. Monitor overhead is within run-to-run variance of the detector alone. 300 ID frames produce zero false transitions (98.3% NOMINAL, brief single-frame breaches contained by the sustained-breach rule); gating scenario replay passes 7/7; log completeness 15/15 checks.

F. Fault injection (Table 6). 16 conditions (clean + 5 corruptions $\times$ 3 severities) on a seeded 300-frame kitti-test subset, report-only. The monitor tracks abrupt degradation: fog:high collapses mAP50 0.847 $\rightarrow$ 0.071 and the monitor flags 99% of frames non-NOMINAL; blur:high and noise:high reach 78%/77% FAIL_SAFE_REQUEST as mAP halves or worse. **The campaign's most important result is negative:** low_light:high removes 19% of mAP50 (0.847 $\rightarrow$ 0.689) while the monitor stays 97.7% NOMINAL; dead_pixels:medium behaves similarly. Mechanism: max-confidence sees the *strongest surviving* detection, not the *missing* ones — gradual recall erosion is invisible to it. All 16 conditions stay within the latency budget (max p95 31.2 ms).

VI. Safety Case and SOTIF Argument

The GSN case (safety/gsn.svg) argues the bounded top claim — *evidence-backed detection of degraded perception conditions and fail-safe requests within the 40 ms budget, within the defined ODD and stated assumptions* — through five sub-goals (hazard-driven requirements; measured baseline; monitor effectiveness; latency; verified traceability), each closed by scripted evidence, with an explicit residual node. The SOTIF argument classifies operating regions: in-ODD operation and abrupt severe degradation are argued acceptably handled (known-safe / known-unsafe-mitigated); gradual low-signal degradation is *known-unsafe-unmitigated* — discovered and quantified by EXP-010, previously unknown-unsafe; unexplored real-world conditions remain unknown-unsafe by definition. Traceability is machine-readable end-to-end: requirements.csv (derivations), traceability_matrix.csv (SR $\rightarrow$ implementation $\rightarrow$ test $\rightarrow$ metric $\rightarrow$ result $\rightarrow$ evidence, all six verified), evidence_index.csv (24 claims $\rightarrow$ artifacts $\rightarrow$ experiments).

VII. Limitations

(1) Ratings in HARA-lite are illustrative, not field-calibrated; one hazard worked fully. (2) The planner/actuator boundary is assumed (A-03), never exercised — fail-safe is a *request*. (3) Corruptions are plausibility models (fog = uniform haze), not weather physics; BDD fog n=13. (4) The low-light blind spot is documented, not fixed; candidate mitigations (detection-count plausibility, temporal consistency, Mahalanobis) were pre-registered scope cuts. (5) Thresholds are KITTI-val-specific; the clear-day control shows they do not transfer across camera domains. (6) INT8 quantization was deferred; FP16 sufficed for the budget. (7) Single dataset family in-ODD; single hardware target.

VIII. Conclusion

A small detector, a simple monitor, and disciplined evidence hygiene are enough to build a genuine, bounded safety argument — including the part most prototypes omit: a quantified statement of where the monitor is blind. The template (hazard analysis → measurable requirements → frozen-threshold runtime monitor → scripted verification → GSN/SOTIF packaging, all seeded and traceable) is the transferable artifact; the low-light result is the demonstration that the template surfaces real residual risk instead of hiding it.

References

1. D. Hendrycks, K. Gimpel, "A Baseline for Detecting Misclassified and Out-of-Distribution Examples in Neural Networks," ICLR 2017.
2. W. Liu et al., "Energy-based Out-of-distribution Detection," NeurIPS 2020.
3. K. Lee et al., "A Simple Unified Framework for Detecting Out-of-Distribution Samples and Adversarial Attacks," NeurIPS 2018.
4. C. Guo et al., "On Calibration of Modern Neural Networks," ICML 2017.
5. N. Leveson, *Engineering a Safer World: Systems Thinking Applied to Safety*, MIT Press, 2011.
6. ISO 21448:2022, "Road vehicles — Safety of the intended functionality."
7. ISO/PAS 8800:2024, "Road vehicles — Safety and artificial intelligence."
8. SCSC Assurance Case Working Group, "Goal Structuring Notation Community Standard," scsc.uk/gsn.
9. D. Hendrycks, T. Dietterich, "Benchmarking Neural Network Robustness to Common Corruptions and Perturbations," ICLR 2019.
10. A. Geiger, P. Lenz, R. Urtasun, "Are we ready for autonomous driving? The KITTI vision benchmark suite," CVPR 2012.
11. F. Yu et al., "BDD100K: A Diverse Driving Dataset for Heterogeneous Multitask Learning," CVPR 2020.
12. G. Jocher et al., "Ultralytics YOLOv8," 2023.